

Lab 14 - CNN

Name: Muskan Goenka | SRN : PES2UG23CS355 | Section: F

1. Introduction

The objective of this lab was to design, implement, and train a Convolutional Neural Network (CNN) capable of classifying hand gesture images into three classes: **rock**, **paper**, and **scissors**. Using PyTorch, the task involved completing a partially implemented Jupyter Notebook, performing data preprocessing, building a CNN model, training it on the Rock-Paper-Scissors dataset, evaluating its accuracy, and finally demonstrating predictions using the trained model.

2. Model Architecture

The CNN model designed for this task consists of **three convolutional blocks** followed by a **fully connected classifier**.

Convolutional Blocks

Each block contains:

- **Conv2d layer**
 - Kernel size: **3 × 3**
 - Padding: **1**
 - Channels:
 - Block 1: $3 \rightarrow 16$
 - Block 2: $16 \rightarrow 32$
 - Block 3: $32 \rightarrow 64$
- **ReLU activation**
- **MaxPool2d(2)** which halves the spatial dimensions

```

... RPS_CNN(
    (conv_block): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (3): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (4): ReLU()
        (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (6): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (7): ReLU()
        (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    )
    (fc): Sequential(
        (0): Flatten(start_dim=1, end_dim=-1)
        (1): Linear(in_features=16384, out_features=256, bias=True)
        (2): ReLU()
        (3): Dropout(p=0.3, inplace=False)
        (4): Linear(in_features=256, out_features=3, bias=True)
    )
)

```

Fully-Connected Classifier

The classifier consists of:

- 1. Flatten layer**

- 2. Linear layer:**

$$64 \times 16 \times 16 \rightarrow 256$$

- 3. ReLU activation**

- 4. Dropout ($p = 0.3$)**

- 5. Final Linear layer:**

$$256 \rightarrow 3 \text{ (for rock, paper, scissors)}$$

This architecture allows the model to progressively learn spatial features and map them into class scores.

3. Training and Performance

Hyperparameters Used

- Optimizer:** Adam
- Learning Rate:** 0.001
- Loss Function:** CrossEntropyLoss
- Batch Size:** 32

- **Epochs:** 10
- **Device:** GPU if available, otherwise CPU

```
Epoch 1/10, Loss = 0.6748
Epoch 2/10, Loss = 0.2257
Epoch 3/10, Loss = 0.1236
Epoch 4/10, Loss = 0.0700
Epoch 5/10, Loss = 0.0487
Epoch 6/10, Loss = 0.0314
Epoch 7/10, Loss = 0.0094
Epoch 8/10, Loss = 0.0054
Epoch 9/10, Loss = 0.0018
Epoch 10/10, Loss = 0.0018
Training complete!
```

Final Test Accuracy

```
# 2. Get the predicted class (highest score)
_, predicted = torch.max(outputs, 1)

total += labels.size(0)
correct += (predicted == labels).sum().item()

print(f"Test Accuracy: {100 * correct / total:.2f}%")

... Test Accuracy: 99.54%
```

Final Test Accuracy: 99.54%

4. Conclusion and Analysis:

This CNN achieved strong performance on the Rock-Paper-Scissors dataset. The model successfully learned visual features of hand gestures and generalized well on the test set. The results indicate that even a relatively lightweight CNN architecture can classify natural images effectively when trained on a clean and structured dataset.

Challenges Faced

- Ensuring correct dataset folder structure in Colab
- Handling image transforms uniformly during training and prediction

- Understanding shape reductions through convolution and pooling layers

Possible Improvements

1. Data Augmentation

Techniques like random rotations, flips, and brightness adjustments would help improve robustness.

2. Deeper Model / More Filters

Adding more convolutional layers or increasing channel counts could improve feature extraction.

3. More Epochs & Learning Rate Scheduling

Training for longer or using schedulers like `StepLR` or `ReduceLROnPlateau` may improve accuracy.

4. Regularization Enhancements

Increasing dropout or using weight decay could prevent overfitting.

```
test_img_path = "/content/dataset/paper/0Uomd0HvQB33m47I.png"
prediction = predict_image(model, test_img_path)
print(f"Model prediction for {test_img_path}: {prediction}")

*** Model prediction for /content/dataset/paper/0Uomd0HvQB33m47I.png: paper
```

```
Randomly selected images:
Image 1: /content/dataset/paper/af6U4DsLPicCZcED.png
Image 2: /content/dataset/scissors/jMVjrtQfvJ0xLT0t.png

Player 1 shows: paper
Player 2 shows: scissors

RESULT: Player 2 wins! scissors beats paper
```